

Finra_wine_reviewers_JintingHang

May 2, 2019

1 Finra Project: Wine reviews

- by Jinting Hang

Questions: 1. find the pair which has the most similar distribution of in items of points 2. compute the word frequency and identify words that differentiate classes the most 3. Build a text classifier using Naive Bayes 4. derive TFIDF matrix of wine reviewers given any taster and further compute similarity matrix 5. identify the taster that has the most similar style of wine reviews

2 Index

- Section [2.1](#)
- Section [3](#)
- Section [3.1](#)
- Section [4](#)
- Section [5](#)

```
In [396]: import numpy as np
import pandas as pd
import re
from itertools import permutations
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.preprocessing import normalize
from scipy.stats import f
from collections import Counter
import matplotlib.pyplot as plt
plt.style.use('ggplot')
%matplotlib inline
```

```
In [397]: df_reviews = pd.read_csv('winemag-data-130k-v2.csv')
df_reviews.head()
```

```
Out[397]:
```

	Unnamed: 0	country	description \
0	0	Italy	Aromas include tropical fruit, broom, brimston...
1	1	Portugal	This is ripe and fruity, a wine that is smooth...
2	2	US	Tart and snappy, the flavors of lime flesh and...

```

3          3          US  Pineapple rind, lemon pith and orange blossom ...
4          4          US  Much like the regular bottling from 2012, this...

```

```

          designation  points  price  province \
0          Vulkà Bianco      87   NaN  Sicily & Sardinia
1          Avidagos        87  15.0          Douro
2          NaN            87  14.0          Oregon
3          Reserve Late Harvest      87  13.0          Michigan
4  Vintner's Reserve Wild Child Block      87  65.0          Oregon

```

```

          region_1      region_2      taster_name \
0          Etna          NaN      Kerin OKeefe
1          NaN          NaN      Roger Voss
2  Willamette Valley  Willamette Valley      Paul Gregutt
3  Lake Michigan Shore          NaN  Alexander Peartree
4  Willamette Valley  Willamette Valley      Paul Gregutt

```

```

taster_twitter_handle          title \
0      @kerinokeefe          Nicosia 2013 Vulkà Bianco (Etna)
1      @vossroger      Quinta dos Avidagos 2011 Avidagos Red (Douro)
2      @paulgwineã      Rainstorm 2013 Pinot Gris (Willamette Valley)
3          NaN  St. Julian 2013 Reserve Late Harvest Riesling ...
4      @paulgwineã  Sweet Cheeks 2012 Vintner's Reserve Wild Child...

```

```

          variety          winery
0      White Blend          Nicosia
1  Portuguese Red  Quinta dos Avidagos
2      Pinot Gris      Rainstorm
3      Riesling      St. Julian
4      Pinot Noir      Sweet Cheeks

```

2.1 Answer question 1

- find the pair which has the moest similar distribution of in items of points

```
In [3]: df_reviews = df_reviews.dropna()
```

```
In [5]: len(df_reviews['variety'].unique())
```

```
Out[5]: 183
```

```
In [69]: df_reviews.groupby(['variety'])['points'].value_counts()
```

```
Out[69]: variety      points
Abouriou          85         1
Aglianico          92         2
                86         1
Albariño           87        13
                88         6
```

	90	6
	89	5
	91	4
	86	3
	92	3
	84	2
	82	1
	85	1
	93	1
Alicante Bouschet	87	1
	88	1
	90	1
	91	1
Aligoté	87	2
Alvarelhão	85	1
Arneis	88	1
	90	1
Auxerrois	90	4
	89	1
Baco Noir	87	1
Barbera	89	23
	90	23
	88	16
	91	14
	86	10
		...
White Blend	83	11
	84	10
	93	5
	94	4
	95	2
	80	1
	82	1
White Riesling	88	1
	96	1
Zinfandel	90	195
	88	138
	91	129
	87	127
	92	107
	89	87
	86	82
	85	65
	93	63
	84	35
	94	25
	83	23
	82	18

	95	8
	80	6
	81	3
	96	2
	97	1
Zweigelt	90	2
	87	1
	88	1

Name: points, Length: 981, dtype: int64

```
In [108]: def similar_distribution_adv(df):
          X= df['points'].value_counts()
          return pd.Series({'Variance': X.var(), 'length':len(X)})
```

```
In [109]: df_variety_distribution = df_reviews.groupby(['variety']).apply(similar_distribution_adv)
```

```
In [120]: df_variety_distribution.dropna(inplace = True)
          df_variety_distribution = df_variety_distribution.loc[df_variety_distribution['Variance'] > 0]
```

```
In [121]: df_variety_distribution.head()
```

```
Out[121]:
```

	Variance	length
variety		
Aglianico	0.500000	2.0
Albariño	12.290909	11.0
Auxerrois	4.500000	2.0
Barbera	69.072727	11.0
Blafränkisch	0.200000	5.0

```
In [393]: alpha = 0.05 # threshold value to test whether two have equal variance.
```


```
In [212]: # Get all permutations of length 2
          # calculate p_value based on ftest to see which pair has the smallest pvalue, normal
          #0.5, we can accept the hypothesis that two groups have the same variance, otherwise
```

```
def most_similar_pair(DF):
    ind = DF.index
    perm = permutations(ind, 2)
    similar_pair = []
    minp = 1 << 30
    for a, b in perm:
        #print('a:',a,'b:', b, DF.loc[a, 'Variance'])
        F = DF.loc[a, 'Variance'] / DF.loc[b, 'Variance']
        df1, df2 = DF.loc[a, 'length'], DF.loc[b, 'length']
        p_value = f.cdf(F, df1, df2)
        if p_value < minp:
            similar_pair = [a, b]
            minp = p_value
    return similar_pair, p_value
```

```
In [215]: similar_pair, pvalue = most_similar_pair(df_variety_distribution)
          similar_pair, pvalue
```

```
Out[215]: (['Fumé Blanc', 'Pinot Noir'], 1.4019283822650899e-06)
```

```
In [216]: # plot out the most similar distribution in in terms of points
def plot_similar_pair(df, similar_pair):
    plt.figure(figsize = (12, 8))
    plt.subplot(211)
    df_test = df.loc[df['variety'] == similar_pair[0], 'points']
    df_test.hist(bins = 50)
    plt.subplot(212)
    df_test = df.loc[df['variety'] == similar_pair[1], 'points']
    df_test.hist(bins = 50)
    plot_similar_pair(df_reviews, similar_pair)
```



2.1.1 We have chosen the pair group which has the smallest pvalue as having the most similar distribution.

3 Answer question 2

- compute the word frequency and identify words that differentiate classes the most

```
In [407]: df_reviews['class'] = 0
          df_reviews.loc[df_reviews['points'] > 90, 'class'] = 1
```

```

In [408]: df_reviews['class'].value_counts()

Out[408]: 0      96336
          1      33635
          Name: class, dtype: int64

In [409]: df_reviews['description'] = df_reviews['description'] + ' '

In [ ]: # analyze the frequency of words for each description
def extract_frequency(df):
    #extract = Counter([re.sub(r'[\w]', ' ', c) for c in df['description'].sum().split(' ')])
    extract = Counter([c.strip(",.() ").replace(" " " ", "") for c in df['description'].sum().split(' ')])
    return pd.Series(extract)

In [ ]: frequency_word = df_reviews.groupby(['class']).apply(extract_frequency).unstack(level=1)

In [ ]: df_reviews_test = df_reviews.iloc[0: 2000]
        frequency_word_test = df_reviews_test.groupby(['class']).apply(extract_frequency).unstack(level=1)

In [ ]: frequency_word_test.iloc[:50]

In [229]: frequency_word.fillna(0, inplace = True)

In [398]: # choose the most frequent 50 words for each class
def distinct_words(df):
    class_0_freq_words = df.sort_values(by = 0, ascending = False).iloc[0:50].index
    class_1_freq_words = df.sort_values(by = 1, ascending = False).iloc[0:50].index
    # use a set intersection to find out the distinct words for each class
    return set(class_0_freq_words) - set(class_1_freq_words), set(class_1_freq_words)

In [400]: distinct_words(frequency_word)

Out[400]: ({'Sauvignon', 'apple', 'has', 'light', 'offers', 'ripe', 'tart', 'vanilla'},
           {'blackberry',
            'bottling',
            'dark',
            'dried',
            'pepper',
            'shows',
            'through',
            'vineyard'})

```

3.1 Answer question 3

- Build a text classifier using Naive Bayes

```

In [ ]: import nltk
        from nltk.corpus import stopwords
        from nltk.stem import WordNetLemmatizer
        import string
        import pandas as pd
        from nltk import pos_tag
        from nltk.stem import PorterStemmer

```

```

In [ ]: def preprocessing(text):
    text2 = " ".join("".join([" " if ch in string.punctuation else ch
for ch in text])).split()
    tokens = [word for sent in nltk.sent_tokenize(text2) for word in
nltk.word_tokenize(sent)]
    tokens = [word.lower() for word in tokens]
    stopwds = stopwords.words('english')
    tokens = [token for token in tokens if token not in stopwds]
    tokens = [word for word in tokens if len(word)>=3]
    stemmer = PorterStemmer()
    tokens = [stemmer.stem(word) for word in tokens]
    tagged_corpus = pos_tag(tokens)
    Noun_tags = ['NN', 'NNP', 'NNPS', 'NNS']
    Verb_tags = ['VB', 'VBD', 'VBG', 'VBN', 'VBP', 'VBZ']
    lemmatizer = WordNetLemmatizer()
    def prat_lemmatize(token, tag):
        if tag in Noun_tags:
            return lemmatizer.lemmatize(token, 'n')
        elif tag in Verb_tags:
            return lemmatizer.lemmatize(token, 'v')
        else:
            return lemmatizer.lemmatize(token, 'n')
    pre_proc_text = " ".join([prat_lemmatize(token, tag) for token, tag in tagged_co
return pre_proc_text

In [238]: ## choose the first 2000 rows data
df_reviews_mini = df_reviews.iloc[:2000]
df_reviews_mini['class'].value_counts()

Out[238]: 0      1171
          1       829
          Name: class, dtype: int64

In [239]: # setting the prior, since 0s and 1s have different probability
p_0 = 1171/2000
p_1 = 829/2000

In [244]: frequency_word_mini = df_reviews_test.groupby(['class']).apply(extract_frequency).un
frequency_word_mini.fillna(0, inplace = True)

In [250]: frequency_word_mini = frequency_word_mini.sort_values(by = 0, ascending = False)
frequency_word_mini.head()

Out[250]: class      0      1
and      3084.0  2723.0
of       1641.0  1503.0
the      1550.0  1385.0
a        1449.0  1169.0
with     957.0   824.0

```

```
In [276]: frequency_word_mini.shape[0]
```

```
Out[276]: 7491
```

```
In [277]: vocabulary = frequency_word_mini.index
vocabulary_class0 = frequency_word_mini.loc[frequency_word_mini[0] != 0].index
vocabulary_class1 = frequency_word_mini.loc[frequency_word_mini[1] != 0].index
```

```
In [279]: len(vocabulary_class0)
```

```
Out[279]: 5131
```

```
In [251]: df_reviews_mini.head()
```

```
Out[251]:
```

	Unnamed: 0	country	description
4	4	US	Much like the regular bottling from 2012, this...
10	10	US	Soft, supple plum envelopes an oaky structure ...
23	23	US	This wine from the Geneseo district offers aro...
25	25	US	Oak and earth intermingle around robust aromas...
35	35	US	As with many of the Erath 2010 vineyard design...

		designation	points	price	province
4	Vintner's Reserve	Wild Child Block	87	65.0	Oregon
10		Mountain Cuvée	87	19.0	California
23		Signature Selection	87	22.0	California
25		King Ridge Vineyard	87	69.0	California
35		Hyland	86	50.0	Oregon

	region_1	region_2	taster_name
4	Willamette Valley	Willamette Valley	Paul Gregutt
10	Napa Valley	Napa	Virginie Boone
23	Paso Robles	Central Coast	Matt Kettmann
25	Sonoma Coast	Sonoma	Virginie Boone
35	McMinnville	Willamette Valley	Paul Gregutt

	taster_twitter_handle	title
4	@paulgwineã	Sweet Cheeks 2012 Vintner's Reserve Wild Child...
10	@vboone	Kirkland Signature 2011 Mountain Cuvée Caberne...
23	@mattkettmann	Bianchi 2011 Signature Selection Merlot (Paso ...
25	@vboone	Castello di Amorosa 2011 King Ridge Vineyard P...
35	@paulgwineã	Erath 2010 Hyland Pinot Noir (McMinnville)

	variety	winery	class
4	Pinot Noir	Sweet Cheeks	0
10	Cabernet Sauvignon	Kirkland Signature	0
23	Merlot	Bianchi	0
25	Pinot Noir	Castello di Amorosa	0
35	Pinot Noir	Erath	0


```

In [349]: X = df_reviews_mini['description']
          y = df_reviews_mini['class']
          Xtrain, Xtest, ytrain, ytest = train_test_split(X, y)
          Xtrain = pd.DataFrame(Xtrain, columns = ['description'])
          Xtrain.index.name = 'ID'
          ytrain = pd.DataFrame(ytrain, columns = ['class'])
          Xtest = pd.DataFrame(Xtest, columns = ['description'])
          Xtest.index.name = 'ID'
          ytest = pd.DataFrame(ytest, columns = ['class'])

In [350]: # since it is a binary classification, we first calculate the likelihood function for
def assign_class(df):
    df_freq = frequency_word_mini
    val = df.values[0][0]
    _dic = Counter([word.strip(',') for word in val.split()])
    # get conditional prob for class 1 and 0
    P_class0, P_class1 = 1, 1
    for key in _dic:
        try:
            P_class1 *= ((df_freq.loc[key, 1] + 1) / (len(vocabulary_class1) + len(vocabulary_class1)))
            P_class0 *= ((df_freq.loc[key, 0] + 1) / (len(vocabulary_class0) + len(vocabulary_class0)))
        except:
            continue

    P_class0 *= p_0
    P_class1 *= p_1
    return pd.Series(0, index = ['Class']) if P_class0 > P_class1 else pd.Series(1, index = ['Class'])

In [351]: y_train_pred = Xtrain.groupby('ID').apply(assign_class)
          y_train_pred['Class'].value_counts()

Out[351]: 0    1101
          1     399
          Name: Class, dtype: int64

In [354]: print('classification_report_train: \n', classification_report(ytrain, y_train_pred))

classification_report_train:
              precision    recall  f1-score   support

0               0.59         0.74         0.65         879
1               0.42         0.27         0.33         621

avg / total           0.52         0.54         0.52        1500

```

```

In [355]: y_test_pred = Xtest.groupby('ID').apply(assign_class)
          y_test_pred['Class'].value_counts()

```

```
Out[355]: 0    372
          1    128
          Name: Class, dtype: int64
```

```
In [356]: print('classification_report_test: \n',classification_report(ytest, y_test_pred))
```

```
classification_report_test:
              precision    recall  f1-score   support

     0       0.60      0.76      0.67      292
     1       0.46      0.28      0.35      208

 avg / total       0.54      0.56      0.54      500
```

The low precision and recall might due to the fact that I haven't removed any stop words and stemmings, if I have more time, I will implement them on both the train and test data.

4 Answer question 4

- derive TFIDF matrix of wine reviewers given any taster and further compute similarity matrix

```
In [357]: df_reviews_mini.head()
```

```
Out[357]: Unnamed: 0 country description \
4          4      US  Much like the regular bottling from 2012, this...
10         10      US  Soft, supple plum envelopes an oaky structure ...
23         23      US  This wine from the Geneseo district offers aro...
25         25      US  Oak and earth intermingle around robust aromas...
35         35      US  As with many of the Erath 2010 vineyard design...

              designation  points  price  province \
4  Vintner's Reserve Wild Child Block      87   65.0    Oregon
10              Mountain Cuvée      87   19.0  California
23              Signature Selection      87   22.0  California
25              King Ridge Vineyard      87   69.0  California
35              Hyland      86   50.0    Oregon

              region_1      region_2  taster_name \
4  Willamette Valley  Willamette Valley  Paul Gregutt
10      Napa Valley      Napa  Virginie Boone
23      Paso Robles  Central Coast  Matt Kettmann
25      Sonoma Coast      Sonoma  Virginie Boone
35      McMinnville  Willamette Valley  Paul Gregutt

taster_twitter_handle title \
```

```

4          @paulgwineã Sweet Cheeks 2012 Vintner's Reserve Wild Child...
10         @vboone Kirkland Signature 2011 Mountain Cuvée Caberne...
23        @mattkettmann Bianchi 2011 Signature Selection Merlot (Paso ...
25         @vboone Castello di Amorosa 2011 King Ridge Vineyard P...
35        @paulgwineã Erath 2010 Hyland Pinot Noir (McMinnville)

```

```

          variety          winery  class
4      Pinot Noir      Sweet Cheeks      0
10  Cabernet Sauvignon  Kirkland Signature      0
23      Merlot          Bianchi      0
25      Pinot Noir  Castello di Amorosa      0
35      Pinot Noir          Erath      0

```

In [379]: # analyze the TFIDF of words for each description

```

def extract_TFIDF(df):
    #extract = Counter([re.sub(r'[\w]', ' ', c) for c in df['description'].sum().sp
    extract = Counter([c.strip(",. ").replace(" " " ", "") for c in df['description
    n = sum(extract.values())
    # calculate the inverse frequency
    dic = {}
    for key in extract:
        dic[key] = np.log(n/ extract[key])

    return pd.Series(dic)

```

In [380]: # IFIDF: inverse frequency of terms

```

grp_by_taster = df_reviews_mini.groupby('taster_name').apply(extract_TFIDF).unstack()
grp_by_taster.head()

```

```

Out[380]: taster_name  Jim Gordon  Joe Czerwinski  Matt Kettmann  Michael Schachner  \
          0.0          0.000000          0.000000          0.0
"dry,"          0.0          5.758902          0.000000          0.0
$12             0.0          0.000000          0.000000          0.0
$20             0.0          0.000000          0.000000          0.0
$25             0.0          0.000000          9.903188          0.0

taster_name  Paul Gregutt  Roger Voss  Sean P. Sullivan  Susan Kostrzewa  \
          0.0000          0.0          9.299907          0.0
"dry,"          0.0000          0.0          0.000000          0.0
$12             0.0000          0.0          0.000000          0.0
$20             9.3299          0.0          0.000000          0.0
$25             0.0000          0.0          0.000000          0.0

taster_name  Virginie Boone
          0.000000
"dry,"          0.000000
$12          10.088846
$20          0.000000
$25          0.000000

```

```
In [385]: # build similarity matrix
A = normalize(grp_by_taster)
similarity_matrix = A.T.dot(A)
np.fill_diagonal(similarity_matrix, 0)
```

```
In [388]: similarity_matrix = pd.DataFrame(similarity_matrix, columns = grp_by_taster.columns,
similarity_matrix
```

```
Out[388]: taster_name      Jim Gordon   Joe Czerwinski   Matt Kettmann  \
taster_name
Jim Gordon      0.000000      17.200131      180.847054
Joe Czerwinski   17.200131      0.000000      22.453870
Matt Kettmann    180.847054      22.453870      0.000000
Michael Schachner 15.393099      6.379032      21.730015
Paul Gregutt     215.180184      27.393243      375.441039
Roger Voss       2.001555      1.396979      2.937266
Sean P. Sullivan 122.863583      16.667886      204.398779
Susan Kostrzewa   9.964144      3.840653      11.617836
Virginie Boone   207.674775      21.176357      404.497218

taster_name      Michael Schachner   Paul Gregutt   Roger Voss  \
taster_name
Jim Gordon      15.393099      215.180184      2.001555
Joe Czerwinski   6.379032      27.393243      1.396979
Matt Kettmann    21.730015      375.441039      2.937266
Michael Schachner 0.000000      25.536940      1.124541
Paul Gregutt     25.536940      0.000000      2.745645
Roger Voss       1.124541      2.745645      0.000000
Sean P. Sullivan 16.742858      248.903781      2.504328
Susan Kostrzewa   3.146881      13.895070      0.915540
Virginie Boone   24.789757      465.042123      2.921999

taster_name      Sean P. Sullivan   Susan Kostrzewa   Virginie Boone
taster_name
Jim Gordon      122.863583      9.964144      207.674775
Joe Czerwinski   16.667886      3.840653      21.176357
Matt Kettmann    204.398779      11.617836      404.497218
Michael Schachner 16.742858      3.146881      24.789757
Paul Gregutt     248.903781      13.895070      465.042123
Roger Voss       2.504328      0.915540      2.921999
Sean P. Sullivan 0.000000      10.069487      239.092529
Susan Kostrzewa   10.069487      0.000000      13.887551
Virginie Boone   239.092529      13.887551      0.000000
```

5 Answer question 5

- identify the taster that has the most similar style of wine reviews

```

In [394]: def find_max(df):
           return df.idxmax()

In [395]: similarity_matrix.apply(find_max, axis = 0)

Out[395]: taster_name
           Jim Gordon      Paul Gregutt
           Joe Czerwinski    Paul Gregutt
           Matt Kettmann    Virginie Boone
           Michael Schachner  Paul Gregutt
           Paul Gregutt      Virginie Boone
           Roger Voss        Matt Kettmann
           Sean P. Sullivan  Paul Gregutt
           Susan Kostrzewa   Paul Gregutt
           Virginie Boone    Paul Gregutt
dtype: object

In [ ]:

```